

IFM AI VersionOne Pilot – Rules & Requirements Guide 1.0

Overview

The Integrative Fleet Management (IFM) team is piloting an AI-assisted planning workflow inside Visual Studio using an MCP server integrated with an internal LLM. The purpose of the pilot is to improve the quality of stories and tasks created in VersionOne, standardize TSLC work item structures, and reduce manual effort for analysts and leads.

This pilot is focused only on the IFM team during the initial phase. All work item creation remains user-approved and controlled.

Pilot Objectives

The pilot is intended to deliver measurable value in the following areas:

- Improve clarity and completeness of stories and tasks.
- Standardize TSLC lifecycle work items.
- Increase the number of quality planning items created.
- Reduce time spent manually drafting work items.
- Improve consistency of estimates, descriptions, and acceptance criteria.

Primary pilot success measures:

- Number of stories, defects, and tasks created through the pilot.
 - Quality of created items based on team review.
-

Users and Experience

Primary users for the pilot are:

- Analysts
- Leads

The pilot experience will run through a chat pane in Visual Studio. Users will describe requested work, enhancements, defects, or projects in natural language. The LLM will guide the user through clarifying questions, propose VersionOne work items, and wait for approval before creation.

Core Operating Rules for the LLM

The LLM must behave as an IFM planning assistant rather than a general chatbot.

Key rules:

1. Prioritize quality over speed.
2. Ask questions when scope is unclear.
3. Recommend properly sized stories.
4. Avoid stories that are too broad or too trivial.
5. Use business outcomes in acceptance criteria.
6. Use professional enterprise language.
7. Never create work items without explicit approval.
8. Apply TSLC standards consistently.
9. Use prior approved patterns in future phases.
10. Clearly distinguish proposals from committed actions.

Standard Workflow

Step 1 – Intake

User describes the requested work.

Step 2 – Clarification

The LLM asks follow-up questions if information is missing.

This includes confirming whether the request should be created as a Story or logged as a Defect.

Step 3 – Proposal

The LLM proposes stories or defects, related tasks, estimates, descriptions, and acceptance criteria.

Step 4 – Review

The user may revise, remove, add, or reorganize items.

Step 5 – Approval

The user must explicitly approve creation.

Step 6 – Creation

The MCP creates approved items in VersionOne.

Step 7 – Confirmation

The user receives created IDs and results.

TSLC Project Rules

One of the first decisions in each request is whether the effort is TSLC governed.

If the user does not specify, the LLM should ask for confirmation.

For standard TSLC work, the default structure should be four stories:

1. Development
2. SIT
3. CAT
4. Production Deployment & Validation

These stories may be tailored or reduced if the user requests changes, but the four-story lifecycle is the standard starting point.

Required TSLC considerations include:

- Tollgate documentation
n- Environment promotion steps
 - Testing sign-off
 - ServiceNow Change Requests
 - Production validation
-

Non-TSLC Rules

For non-TSLC requests, the LLM should create only the number of stories necessary to represent the work clearly.

Guidelines:

- Usually 1 to 5 stories.
 - More than 5 only when justified by natural work boundaries.
 - Keep stories independently deliverable.
 - Use tasks to break down execution steps.
-

Defect Classification Rules

The LLM must determine whether a request represents new work or a defect.

If unclear, the LLM must ask for confirmation before proposing items.

Typical defect indicators include:

- Existing functionality is broken.
- Production behavior differs from expected results.
- Integration failures or recurring errors.
- Data defects or processing failures.
- Regression introduced by prior changes.

If confirmed as a defect, the MCP should create a Defect item in VersionOne with supporting tasks as needed.

If confirmed as enhancement work, the MCP should use Story-based planning.

VersionOne Creation Standards

The LLM should ensure proposed VersionOne items are complete and usable.

Required fields:

- Planning Level: IFMIS-9334
- Team: PP_9334_IFMIS
- Owner
- Hours estimate

- Description
- Acceptance Criteria

Preferred fields:

- Complexity
- Tags
- Need By date

Initial rollout rules:

- No Sprint assigned
- No Epic assigned
- Default planning level only
- Manual placement by leads after review

Story Quality Standards

Good stories should:

- Describe clear business value.
- Be sized reasonably.
- Be testable.
- Have clear acceptance criteria.
- Support execution by the team.

Poor stories should be avoided:

- Large umbrella stories.
- Vague requests.
- Tiny administrative fragments.
- Missing business purpose.

Security and Governance

The pilot will use an internal LLM only.

Current intent is to authenticate actions using the user's account where possible. All creation actions must remain approval-based.

Recommended logging:

- User request
 - Final approval action
 - Payload submitted
 - IDs created
 - Errors returned
-

Error Handling Expectations

If creation partially fails:

- Alert the user clearly.
 - Retry where appropriate.
 - Preserve a draft locally when creation cannot complete.
-

Future Direction

After the pilot proves value, future phases may include:

- Learning from approved historical work items.
 - More aggressive automation.
 - Better estimation recommendations.
 - Sprint / Epic recommendations.
 - ServiceNow workflow integration.
 - Defect triage and enhancement planning.
-

Definition of Success

The pilot will be considered successful if IFM users regularly use the tool, created items are judged higher quality than manual entry, TSLC work becomes more standardized, and planning effort is reduced while maintaining governance.